



Université
de Lille

AUTOMATIC IMAGE SEGMENTATION BY HISTOGRAM ANALYSIS

RECONNAISSANCE DES FORMES

Iacob Sergiu

3rd March 2020

Contents

| | | |
|----------|--|-----------|
| 1 | Introduction | 2 |
| 1.1 | Context | 2 |
| 1.2 | Motivation | 2 |
| 1.3 | Goals | 2 |
| 2 | Binarisation using Bayes' Theorem | 3 |
| 2.1 | Loading data | 3 |
| 2.2 | Technical acknowledgements | 4 |
| 2.3 | Fixed threshold | 4 |
| 2.4 | A priori class probabilities | 5 |
| 2.5 | Conditional probabilities | 6 |
| 2.6 | Automatic thresholding using Bayes | 7 |
| 2.7 | Segmenting digits | 8 |
| 2.8 | Classification error | 11 |
| 3 | Conclusion | 12 |

CHAPTER 1

INTRODUCTION

1.1 CONTEXT

This report is the result of a university assignment. It aims to prove that the student understood the motivation, goals and means of studying pattern recognition. Therefore, this is a way of summarizing a series of observations and experiments done on images containing shapes.

1.2 MOTIVATION

In order to be able to work with shapes, we must first acquire them. Most of the time, they are not given to us in their simplest form, but found in images with all kinds of backgrounds. They must be extracted as precisely as possible, making sure we differentiate them from the background as accurately as we can.

1.3 GOALS

We'll try to apply *binary segmentation* to a series of gray images. Our goal is to find an automatic method of doing this and identify objects and backgrounds, mainly focusing on the histograms of gray values.

CHAPTER 2

BINARISATION USING BAYES' THEOREM

2.1 LOADING DATA

The image we'll try to apply binarisation to is the following:




FIGURE 2.1
Image to experiment on

Which is composed of the following 2 images:



FIGURE 2.2
Separate images

The images will be loaded as gray images:

```
1 # Chargement d'une image en niveaux de gris
2 rdfReadGreyImage <- function (nom) {
3   image <- readImage (nom)
4   if (length (dim (image)) == 2) {
5     image
6   } else {
7     channel (image, 'red')
8   }
9 }
```

LISTING 2.1
Loading data in R

2.2 TECHNICAL ACKNOWLEDGEMENTS

Please note that for the following scripts, some data was pre-loaded to avoid loading it every time.

```

1  # loading some data only once
2  image <- rdfReadGreyImage ("2classes_100_100_8bits_2016.png")
3  above <- rdfReadGreyImage ("2classes_100_100_8bits_omega1_2016.png")
4  below <- rdfReadGreyImage ("2classes_100_100_8bits_omega2_2016.png")
5
6  nbins <- 256
7  h <- hist (as.vector (image), freq=FALSE, breaks = seq (0, 1, 1 / nbins))
8  h1 <- hist (as.vector (above), freq=FALSE, breaks = seq (0, 1, 1 / nbins))
9  h2 <- hist (as.vector (below), freq=FALSE, breaks = seq (0, 1, 1 / nbins))

```

LISTING 2.2

Loading some data only once

2.3 FIXED THRESHOLD

Let's take a look at the histograms below. Left is for the “merged” image, right is for the separate images, the histograms being plotted together. Based on the left histogram, we can have a “manual” attempt at binarising the image: using a *fixed threshold*. We simply choose a value for the threshold and we assign, for each pixel, one of the predicted classes as it follows:

$$\hat{w}_1 = \{P \in I | I(P) < \hat{X}\}$$

$$\hat{w}_2 = \{P \in I | I(P) \geq \hat{X}\}$$

where \hat{X} is our threshold.

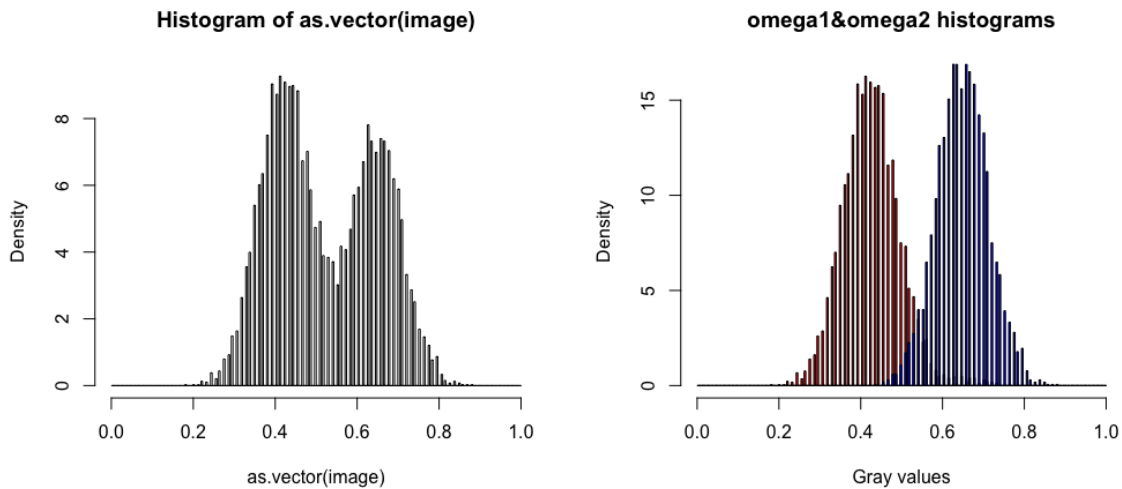


FIGURE 2.3

Histogram of grey values for

```

1  # for each threshold, calculate the binary image
2  # image has values between [0, 1] so our threshold has to be in the same range
3  # for each pixel, the expression (pixel_value - threshold) >= 0 will assign it
   to one of two classes: w1 or w2
4  binaire50 <- (image - 0.5) >= 0
5  display (binaire50, "image binaire 0.35", method="raster", all=TRUE)
6  binaire55 <- (image - 0.55) >= 0
7  display (binaire55, "image binaire 0.35", method="raster", all=TRUE)
8  binaire60 <- (image - 0.6) >= 0
9  display (binaire60, "image binaire 0.35", method="raster", all=TRUE)

```

LISTING 2.3

Using fixed threshold

Let's take a look at the results, using threshold values of 0.5, 0.55 and 0.6.

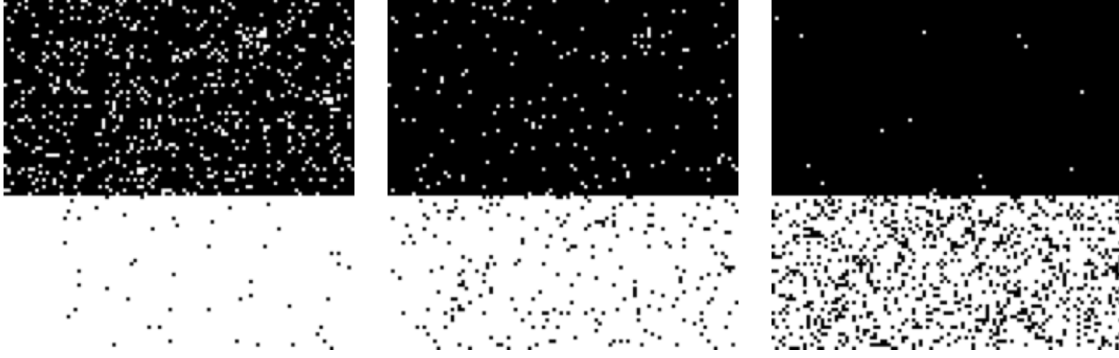


FIGURE 2.4
Fixed threshold results

It is not hard to see that there is a lot of noise in our predictions and that they are far from perfect. The threshold value of 0.6 might seem attractive, but it only manages to classify well most of the pixels from the image on the top. Further more, it heavily missclassifies the pixels from the image on the bottom.

2.4 A PRIORI CLASS PROBABILITIES

Our image 2.1 is formed of 2 separate images 2.2. We can try to calculate the following probabilities:

$P(w_1)$, the probability of a pixel to be in w_1 class

$P(w_2)$, the probability of a pixel to be in w_2 class

Knowing that the dimension of the merged image is 100x100 and the 2 images are 100x43, 100x57, the probabilities above are easy to calculate.

```

1  calculateClassAPrioriProbabilities <- function() {
2      # pw1 = (number_of_pixels_from_the_first_image) / (number_of_total_pixels)
3      # pw1 = (number_of_pixels_from_the_second_image) / (number_of_total_pixels)
4      noTotalPixels = dim(image)[1] * dim(image)[2]
5      pw1 = (dim(above)[1] * dim(above)[2]) / noTotalPixels
6      pw2 = (dim(below)[1] * dim(below)[2]) / noTotalPixels
7      c(pw1, pw2)
8  }
```

LISTING 2.4
A priori class probabilities

In this case, the probabilities are $P(w_1) = 0.57$ and $P(w_2) = 0.43$.

2.5 CONDITIONAL PROBABILITIES

Having $P(w_1)$ and $P(w_2)$, we can take a step further and calculate the following conditional probabilities:

$P(X|w_1)$, the probability of a pixel having the gray value of X if it's in the w_1 class

$P(X|w_2)$, the probability of a pixel having the gray value of X if it's in the w_2 class

These are, according to Wikipedia 2020, equal to:

$$P(X|w) = \frac{P(w|X) * P(X)}{P(w)}$$

We already have $P(w)$ and $P(X)$ is simply the number of pixels with the gray value of X divided by the total number of pixels. $P(w|X)$ is the number of pixels with the value X in the merged image corresponding to w divided by the total number of pixels with a value of X .

```

1  calculateConditionalProbability <- function (X, probs){
2      # let's take a look at the histogram again
3      # we can notice that h$counts[X + 1] gives us the number of pixels that have
4      the gray value of X
5      # the "+ 1" above is because in R arrays start from 1
6
7      # so, P(X | I) = the probability of a pixel from the merged image to have
8      the gray value of X
9      px <- h$counts[X + 1] / sum(h$counts)
10     # P(X | w) = the probability of a pixel from the merged image to have the
11     gray value of X
12     # IF that pixel is from the w class
13     # Using bayes theorem, we get:
14     # P(X|w) = (P(w | X) * P(X)) / P(w)
15     # and P(w | X) = how many pixels of value X are in the image corresponding
16     to w / the total number of pixels with a value of X
17
18     p_w1_if_x <- h1$counts[X + 1] / h$counts[X + 1]
19     p_x_and_w1 <- p_w1_if_x * px
20     p_x_if_w1 <- p_x_and_w1 / probs[1]
21
22     p_w2_if_x <- h2$counts[X + 1] / h$counts[X + 1]
23     p_x_and_w2 <- p_w2_if_x * px
24     p_x_if_w2 <- p_x_and_w2 / probs[2]
25
26     c(px, p_x_if_w1, p_x_if_w2)
27 }
28
29 print (calculateConditionalProbability(141, probs))

```

LISTING 2.5
Calculating conditional probabilities

Running the above function for $X = 141$, we get the following probabilities:

$$P(X|I) = 0.011800000$$

$$P(X|w_1) = 0.008947368$$

$$P(X|w_2) = 0.015581395$$

where $P(X|I)$ refers to the probability of a pixel to have the gray value of X in the whole (merged) image. Of course, $P(X|w_1) * P(w_1) + P(X|w_2) * P(w_2) = P(X|I)$.

2.6 AUTOMATIC THRESHOLDING USING BAYES

Having a threshold \hat{X} , we can define an assignment error as being the following:

$$P(\epsilon|\hat{X}) = \underset{x}{\operatorname{argmin}} \left(\sum_{X \in \hat{w}_2} P(X|w_1) * P(w_1) + \sum_{X \in \hat{w}_1} P(X|w_2) * P(w_2) \right)$$

Our goal is to find the \hat{X} that minimizes $P(\epsilon|\hat{X})$. For this, we can try each value for \hat{X} and see which is the best.

```

1  automaticSegmentationUsingBayes <- function() {
2      # calculate these probabilities only once
3      probs <- calculateClassAPrioriProbabilities()
4      # only get the imageData, it will be faster for comparisons
5      image <- imageData(image)
6      # build our reference segmentation
7      perfect <- matrix (nrow=dim(image)[1], ncol=dim(image)[2])
8      # dim(image)[1] = number of columns
9      for (i in 1:dim(image)[1])
10         for (j in 1:dim(image)[2]){
11             perfect[i, j] <- FALSE
12         }
13     for (i in 1:dim(image)[1])
14         for (j in dim(image)[2]:dim(image)[2]){
15             perfect[i, j] <- TRUE
16         }
17     display(perfect, method="raster", all=TRUE)
18
19     # precalculate all probabilities
20     condProbs <- matrix(nrow = 256, ncol = 2)
21     for (X in 0:255){
22         res <- calculateConditionalProbability(X, probs)
23         condProbs[X + 1, 1] <- res[2]
24         condProbs[X + 1, 2] <- res[3]
25     }
26
27     min_error = 2 * dim(image)[1] * dim(image)[2]
28     for (X in 0:255){
29         binary <- (image - X/255) >= 0
30         # only get the data, otherwise the comparisons will be really slow
31         binary <- imageData(binary)
32         error <- 0
33         # these pixels should be in w1 class
34         for (i in 1:dim(image)[1])
35             for (j in 1:dim(image)[2])
36                 if (perfect[i, j] != binary[i, j]){
37                     # if it should be in the first class
38                     if (j < dim(image)[2])
39                         error <- error + condProbs[image[i, j] * 255 + 1, 1] * probs[1]
40                     else
41                         error <- error + condProbs[image[i, j] * 255 + 1, 2] * probs[2]
42                 }
43         if (error < min_error){
44             min_error <- error
45             best_threshold <- X
46         }
47     }
48     print (c(min_error, best_threshold))
49     # segment using best_threshold
50     binary <- (image - best_threshold/255) >= 0
51     display(binary, method="raster", all=TRUE)
52 }

```

LISTING 2.6
Automatic segmentation using Bayes

Below we can look at the “perfect” image and the one we got using this automatic segmentation. To calculate the “perfect” segmentation, we simply assigned all pixels from the first (above) image to a class, and the rest to another class. According to the script, the minimum error is 1.636 and the best threshold for the gray value is $X = 139$.



FIGURE 2.5
Perfect segmentation and our best result

2.7 SEGMENTING DIGITS

We'll take a look now at the results of this method for some digits. We'll be working with the following images:




FIGURE 2.6
Digit images

For the 0 digit, we have 3 images: one we need to segment and 2 already “segmented” photos. For the second and third segment, we can see that some pixels are black. Those pixels correspond to the w_1 and w_2 classes. We'll be using this fact for the calculation of probabilities.

Let us first load our new new data and work with the 0 digit:

```
1 # get img, w1 and w2 histograms
2 image <- rdfReadGreyImage("rdf-chiffre-0-8bits.png")
3 img_w1 <- rdfReadGreyImage("rdf-chiffre-0-8bits_omega1.png")
4 img_w2 <- rdfReadGreyImage("rdf-chiffre-0-8bits_omega2.png")
5 nbins <- 256
6 h <- hist(as.vector(image), freq=FALSE, breaks = seq(0, 1, 1 / nbins))
7 h1 <- hist(as.vector(img_w1), freq=FALSE, breaks = seq(0, 1, 1 / nbins))
8 h2 <- hist(as.vector(img_w2), freq=FALSE, breaks = seq(0, 1, 1 / nbins))
```

LISTING 2.7
Loading data for digits

To use the automatic bayes thresholding, we need to calculate $P(w)$, $P(X)$ and $P(X|w)$, so $P(w|X)$ as well. $P(w)$ will be equal to the number of pixels from the w class (so the black pixels in the second and third images) divided by the total number of pixels. For $P(X)$, we simply divide the number of pixels with the gray value of X to the total number of pixels.

$P(w|X)$ isn't that tricky. If, for example, the pixels of value X belong to the w_1 class, then in the histogram for w_1 we won't have any X values. That's because those pixels with the value of X were given the 0 value (black pixels).

So, $\frac{\text{count}_{\text{total}} - \text{count}_{w_1}}{\text{count}_{\text{total}}}$ will give us the proportion of pixels of value X that were assigned to the w_1 class. We can use this to calculate $P(w|X)$. Afterwards, we can easily calculate $P(X|w)$ using Bayes' theorem.

```

1  classify0Digit <- function(){
2      # calculate a priori probabilities
3      # P(w) = probability of a pixel to be in the w class
4      # the pixels that are in w class have a gray value of 0 in the image
   corresponding to w
5      probs <- c(h1$counts[1] / sum(h$counts), h2$counts[1] / sum(h$counts))
6
7      # only get the imageData, it will be faster for comparisons
8      image <- imageData(image)
9      img_w1 <- imageData(img_w1)
10     img_w2 <- imageData(img_w2)
11
12     # build our reference segmentation
13     perfect <- matrix(nrow=dim(image)[1], ncol=dim(image)[2])
14     # dim(image)[1] = number of columns
15     for (i in 1:dim(image)[1])
16         for (j in 1:dim(image)[2]){
17             if (img_w1[i, j] == 0){
18                 perfect[i, j] <- TRUE
19             }
20             else{
21                 perfect[i, j] <- FALSE
22             }
23         }
24     display(perfect, method="raster", all=TRUE)
25
26     # precalculate all conditional probabilities
27     condProbs <- matrix(nrow = 256, ncol = 2)
28     for (X in 0:255){
29         px <- h$counts[X + 1] / sum(h$counts)
30         # P(X | w) = probability of a pixel having the gray value of X if the
   pixel is from class w
31         # Using bayes theorem, we get:
32         # P(X|w) = (P(w | X) * P(X)) / P(w)
33         # P(w | X) = the prob. of the pixel being from the w class, if its value
   is X
34         #
35
36         # h1$counts[X + 1] + h2$counts[X + 1] = h$counts[X + 1]
37         # if h1$counts[X + 1] - h$counts[X + 1] is different than 0, then those
   pixels were NOT assigned to w1
38
39         if (px == 0){
40             condProbs[X + 1, 1] <- 0
41             condProbs[X + 1, 2] <- 0
42             next
43         }
44         p_w1_if_x <- (h$counts[X + 1] - h1$counts[X + 1]) / h$counts[X + 1]
45         p_x_and_w1 <- p_w1_if_x * px
46         p_x_if_w1 <- p_x_and_w1 / probs[1]
47         p_w2_if_x <- (h$counts[X + 1] - h2$counts[X + 1]) / h$counts[X + 1]

```

```

48     p_x_and_w2 <- p_w2_if_x * px
49     p_x_if_w2 <- p_x_and_w2 / probs[2]
50
51     condProbs[X + 1, 1] <- p_x_if_w1
52     condProbs[X + 1, 2] <- p_x_if_w2
53 }
54
55 probSum <- 0
56 for (X in 0:255){
57     if(is.nan(condProbs[X + 1, 1]) == FALSE)
58         probSum <- probSum + probs[1] * condProbs[X + 1, 1] + probs[2] *
condProbs[X + 1, 2]
59 }
60
61 print ('Making sure that the sum of conditional probabilities is 1:')
62 print (probSum)
63
64 # search for the best threshold
65 min_error <- 2 * dim(image)[1] * dim(image)[2]
66 best_threshold <- 0
67 for (X in 0:255){
68     binary <- (image - X/255) >= 0
69     # only get the data, otherwise the comparisons will be really slow
70     binary <- imageData(binary)
71
72     error <- 0
73     for (i in 1:dim(image)[1])
74         for (j in 1:dim(image)[2])
75             if (perfect [i, j] != binary [i, j]){
76                 # if it should have been in w1
77                 if (img_w1[i, j] == 0)
78                     error <- error + condProbs[image[i, j] * 255 + 1, 1] * probs[1]
79                 else
80                     # it should have been in w2
81                     error <- error + condProbs[image[i, j] * 255 + 1, 2] * probs[2]
82             }
83
84     if (error < min_error){
85         min_error <- error
86         best_threshold <- X
87     }
88 }
89 print (c (min_error, best_threshold))
90 # segment using best_threshold
91 binary <- (image - best_threshold / 255) >= 0
92 display(binary, method="raster", all=TRUE)
93 }

```

LISTING 2.8
Automatic segmentation for digit 0

Running the above script will give us an error of 0.04583333, the best threshold being $X = 142$. We also applied the same threshold for digit 1 and we can see that it is quite a good threshold. The segmentation isn't perfect, but it gets the job done.



FIGURE 2.7
Automatic segmentation using best threshold for digit 0



FIGURE 2.8
Perfect segmentation

2.8 CLASSIFICATION ERROR

As a final conclusion, we can look at the error rate of our algorithm.

```

1  binary <- (image - best_threshold / 255) >= 0
2  # build our reference segmentation
3  perfect <- matrix (nrow=dim(image)[1], ncol=dim(image)[2])
4  # dim(image)[1] = number of columns
5  for (i in 1:dim(image)[1])
6    for (j in 1:dim(image)[2]){
7      if (img_wl[i, j] == 0){
8        perfect[i, j] <- TRUE
9      }
10     else{
11       perfect[i, j] <- FALSE
12     }
13   }
14  sprintf("Error for digit 0: %f", sum(abs(perfect - binary)) / (dim(perfect)[1] *
15    dim(perfect)[2]))
16
17  # calculate error rate for 1
18  image <- rdfReadGreyImage("rdf-chiffre-1-8bits.png")
19  binary <- (image - best_threshold / 255) < 0
20  perfect <- rdfReadGreyImage("rdf-chiffre-1-8bits_classe_a_trouver.png")
21  display(binary, method="raster", all=TRUE)
22  sprintf("Error for digit 1: %f", sum(abs(perfect - binary)) / (dim(perfect)[1] *
23    dim(perfect)[2]))

```

LISTING 2.9
Calculating error rate

This gave us an error rate of 0.033333, that is around 3.3% misclassified pixels for digit 0. For digit 1, we have an error of 0.029167, so around 3% misclassified pixels.

CHAPTER 3

CONCLUSION

Bayes' theorem relieves us from having to manually search for the best threshold. It provided results close to what we were looking for, with low error rates. It can be used in a multitude of scenarios as a strong classification tool, this report presenting just one simple example.

However, when segmenting the digits, for example, we did rely on already existing perfect segmentations in order to calculate probabilities like $P(w)$. Without those, Bayes' theorem could not have been applied. This is, however, a step in the right direction, especially coming from a method where we were manually selecting the threshold based on the histograms of gray values.

BIBLIOGRAPHY

Wikipedia (2020). *Bayes' Theorem*. Wikimedia Foundation. URL: https://en.wikipedia.org/wiki/Bayes%27_theorem (visited on 27th Feb. 2020).